

---

# Parallel Computing

Linh Nguyen '16

# A (very) Short Introduction to CS

---

GTX 680

- Use computers to solve problems
  - Is applicable in many different disciplines
  - Has many areas
  - Is fun
- 
- NOT about using softwares (such as Word or Excel)

# Why Parallel?

---

GTX 680

Single-threaded performance has plateaued

# Why Parallel?

---

GTX 680

# Terminology

---

- Central Processing Unit = CPUs
- Graphics Processing Unit = GPUs
- Threads = processors that can do arithmetic computations.

# A survey of parallel hardware

---

GTX 680

## Personal Mobile Device

2 CPU cores/  
3 GPU cores

4 CPU cores/  
4 GPU cores

iPhone 5

Galaxy S3

# A survey of parallel hardware

---

GTX 680

## Desktop Space

# A survey of parallel hardware

---

GTX 680

## Warehouse Space

2012



# A survey of parallel hardware

GTX 680

## Warehouse Space

- ▮ 2508 CPU cores
- ▮ 187264 GPU cores

# CPU vs GPU

GTX 680

## What is a CPU?

- CPU
  - SR71 Jet
- Capacity
  - 2 passengers
- Top Speed
  - 2200 mph

# CPU vs GPU

GTX 680

## What is a GPU?

- GPU
  - Boeing 747
- Capacity
  - 605 passengers
- Top Speed
  - 570 mph

# CPU vs GPU


GTX 680

	Capacity (passengers)	Speed (mph)	Throughput (passengers * mph)
"CPU" Fighter Jet	2	2200	4400
"GPU" 747	452	555	250,860

- 
- Assign yourself the number 1
  - Pair off with someone standing
  - Add your numbers together and adopt the sum as your new number
  - One of the pair sits down
  - Repeat

# Parallel Reduction

- For a large array
  - Each thread adds a pair of numbers
  - Write partial sums to the temporary array
  - Repeat until done



# The BIG idea behind CUDA

- Replace loops with a functions (a **kernel**) excecuting at each point in a problem domain
  - E.g., process a 1024x1024 image with one kernel invocation per pixel or  $1024 \times 1024 = 1,048,576$  kernel executions.

## Traditional loops    Data Parallel CUDA

```
void
vecAdd(const int n,
       const float *a,
       const float *b,
       float *c)
{
    int i;
    for (i = 0; i < n; i++)
        c[i] = a[i] * b[i];
}
```

```
__global__ void
vecAdd(const float *a, const float *b,
       float *c)
{
    int id = threadIdx.x + \
            blockDim.x * blockIdx.x;
    c[id] = a[id] * b[id];
}
// many instances of the kernel,
// called threads, execute
// in parallel
```

# What is GPGPU?

---

- General Purpose computation using GPU in applications (other than 3D graphics)
  - GPU accelerates critical path of application
- Data parallel algorithms leverage GPU attributes
  - Large data arrays, streaming throughput
  - Low-latency floating point (FP) computation
- Applications - see [//GPGPU.org](http://GPGPU.org)
  - Game effects (FX) physics, image processing
  - Physical modeling, computational engineering, matrix algebra, convolution, correlation, sorting





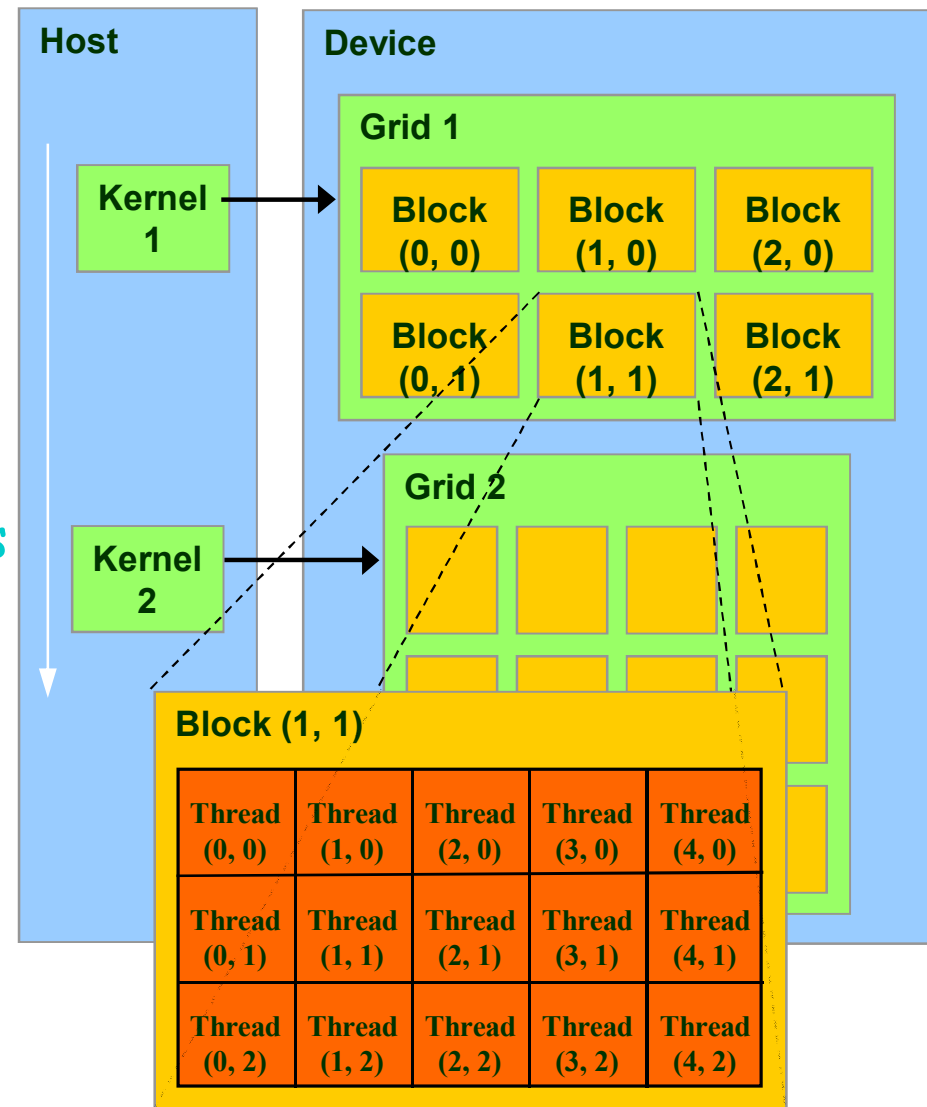
# CUDA Programming Model

---

- The GPU is viewed as a compute **device** that:
  - Is a coprocessor to the CPU or **host**
  - Has its own DRAM (**device memory**)
  - Runs many **threads in parallel**
    - Hardware switching between threads (in 1 cycle)
- Data-parallel portions of an application are executed on the device as **kernels** which run in parallel on many threads

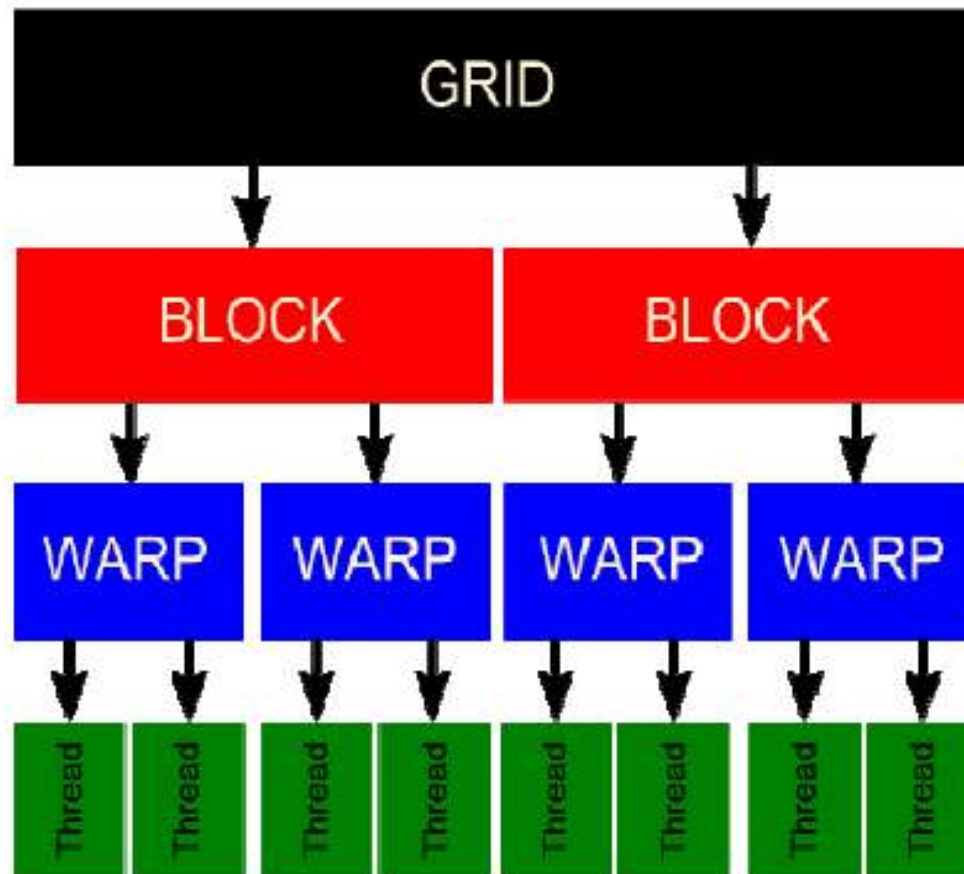
# Thread Scheduling: Grids and Blocks

- Kernel executed as a **grid of thread blocks**
  - All threads share data memory space
- **Thread block** is a batch of threads, can **cooperate** with each other.
- In each thread block there are **warps** of 32 threads.
- Threads and blocks have IDs



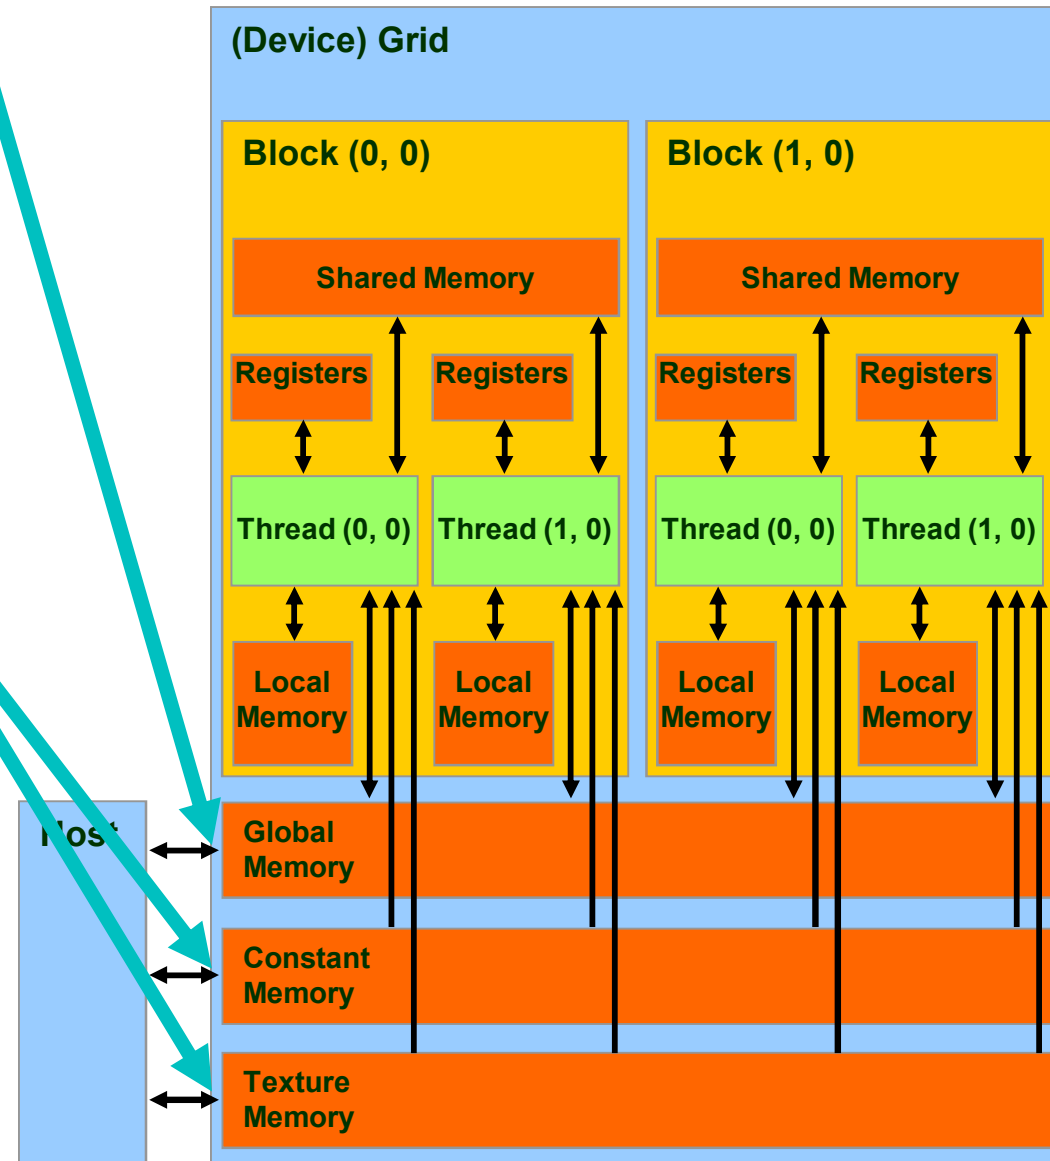
# Execution model

---



# Memory Hierarchy

- Global memory
  - Main means of communicating R/W Data between **host** and **device**
  - Contents visible to all threads
- Texture and Constant Memories
  - Constants initialized by host
  - Contents visible to all threads



Courtesy: NDVIA

# HotSpot

---

- Based on a well-known duality between electric current and heat flow
- Construct a network of thermal “resistances” and “capacitances”
- Solve standard RC circuits with finite element analysis

# The circuit

---

- Divide the chip into small “blocks”
- Each node has a capacitor to model transient state
- The stencil equation

$$\left( \frac{(T_N + T_S - 2T_{i,j,k})}{R_x} + \frac{(T_E + T_W - 2T_{i,j,k})}{R_y} + \frac{(T_A + T_B - 2T_{i,j,k})}{R_z} + P_{i,j,k} \right) \times \frac{\Delta t}{C} = \Delta T$$

# Profiling HotSpot


---

- 512x512 thermal grid
- Intel Xeon(R) X5550 2.67GHz
- Gprof 2.22

# Accelerating Slope Function

---

- Each thread sweeps points in the  $z$  direction, calculating the temperatures using the 7-point stencil
- $X$  and  $Y$  dimensions are blocked with  $A \times B$  thread blocks
- $A$  and  $B$  are user-configurable parameters





# Testing the slope kernel

---

- CUDA toolkit 5.5
- Geforce GT 630M, Tesla C2050 and Tesla K20c

# Testing the HotSpot Program

---

GPU Solver

CPU Solver

# Testing the HotSpot Program

---

